




PostgreSQL



OPEN SOURCE
DEPARTMENT



Day 1 Contents



- Why PostgreSQL?
- History of PostgreSQL.
- ORDBMS concepts and terminology
- PostgreSQL application.
- Columns Data Types
- DML (Insert / update / delete /Truncate)

Why PostgreSQL?



- Designed for high volume environments.
- Cross platform
- Low / No Cost.
- Stability
- Open Source

https://en.wikipedia.org/wiki/PostgreSQL#Prominent_users

<http://www.postgresql.org/about/users/>

<http://www.postgresql.org/about/sponsors/>

https://en.wikipedia.org/wiki/Comparison_of_relational_database_management_systems

History of PostgreSQL



- PostgreSQL is derived from the POSTGRES package written at the University of California at Berkeley by a computer science professor named [Michael Stonebraker](#)
- POSTGRES used PostQUEL as query language



https://en.wikipedia.org/wiki/Michael_Stonebraker

History of PostgreSQL



- In 1994, Berkeley graduate students [Andrew Yu](#) and [Jolly Chen](#) replaced the PostQUEL query language interpreter with one for the SQL query language, creating [Postgres95](#).
- We chose a new name, PostgreSQL, to reflect the relationship between the original POSTGRES and the more recent versions with SQL capability.

Relational database concepts



- Tables

- A Collection of related data. The table has a name; a number of columns and a number of rows, a table in a database looks like a simple spreadsheet.

- Columns

- Each column in the table has a unique name and contains different data. Additionally, each column has an **associated data type** as an integer, strings or Timestamp and so on . Columns are sometimes called fields or attributes.

Relational database concepts



- Rows

- are a group of related data. Because of the tabular format, each row has the same attributes.

- Rows are also called records or tuples.

- Values

- Each row consists of a set of individual values that correspond to columns. Each value must have the data type specified by its column.

Relational database concepts



- Primary Key

- A primary key is unique. A key value can not occur twice in one table. With a key, you can find at most one row.

- Foreign Key

- A foreign key is the linking pin between two tables.

- Referential Integrity

- Referential Integrity makes sure that a foreign key value always points to an existing row.

Relational database concepts



- Index

- it is a data structure that improves the speed of data retrieval operations on a database table
- Disadvantages: Storage Size & Insertion Time

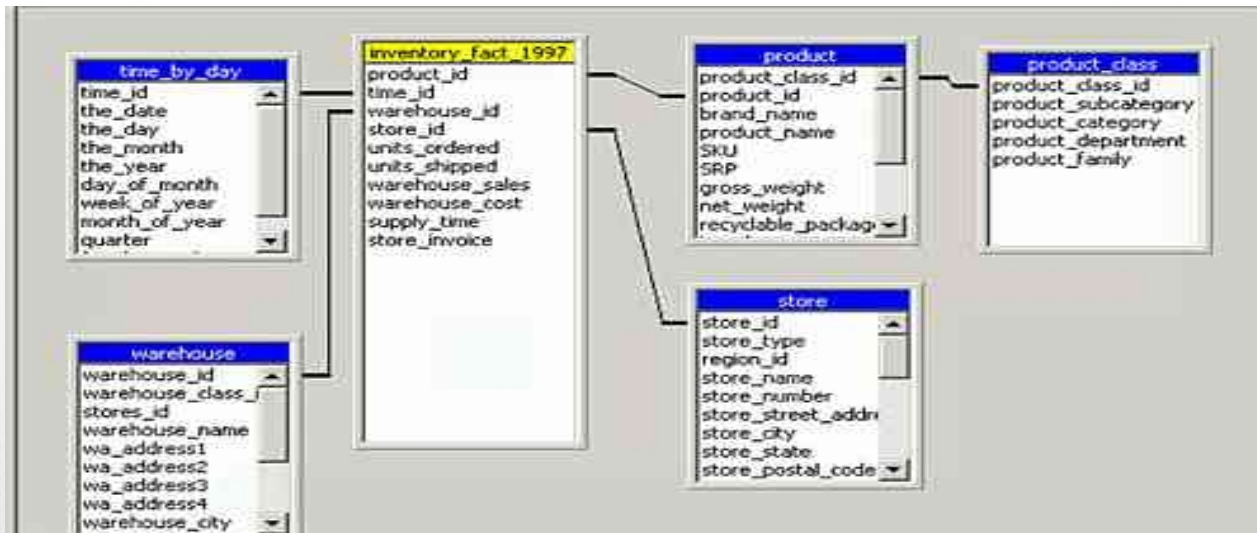
Relational database concepts



OPEN SOURCE
DEPARTMENT

- Schemas

— It is akin to a blueprint for the database. A schema should show the tables along with their columns, and the primary key of each table and any foreign keys. A schema does not include any data.



Relational database



- RDBMS software:
 - Enables you to implement a database with tables, columns and indexes.
 - Guarantees the Referential Integrity between rows of various tables.
 - Interprets an SQL query and combines information from various tables.
 - C/C++, Java Interface





- O for Object:

- The basic goal for the Object-relational database is to bridge the gap between relational databases and the object-oriented modeling used in programming languages such as Java, C++.

- R for Relational:

- It's called relational because all the data is stored into different tables established using primary keys or other keys known as foreign keys.

Report Time!

?

DBMSs Types?



Entity Relationship Diagram

ERD Symbols and Notations



- Entity

- An entity can be a person, place, event, or object that is relevant to a given system. For example, a school system may include students, teachers, major courses, subjects, fees, and other items. Entities are represented in ER diagrams by a rectangle and named using *singular* nouns.

ERD Symbols and Notations



- Attribute

- An attribute is a property or characteristic of an entity, relationship. For example, the attribute Inventory Item Name is an attribute of the entity Inventory Item.

- Multivalued Attribute

- If an attribute can have more than one value it is called an multivalued attribute. For example a person entity can have multiple hobbies values.

ERD Symbols and Notations



- Derived Attribute

An attribute based on another attribute. This is found rarely in ER diagrams. Such as calculations
SUCH AS call duration, age.

- Relationship

A relationship describes how entities interact.

Example: Student, address, track, staff, courses, desk

Designing Your Database



- Simple tables

- that describe a real-world object. They might also contain keys to other simple objects with which they have a one-to-one or one-to-many relation-ship.

- Linking tables

- that describe a many-to-many relationship between two real objects.

Installing PostgreSQL



- For CentOS/Red Hat Distros:

```
# yum install postgresql postgresql-server pgadmin3
```

```
# service postgresql initdb
```

```
# chkconfig postgresql on
```

```
# service postgresql start
```

- For Ubuntu/Debian :

```
$ sudo apt-get install postgresql postgresql-client pgadmin3
```

Postgres Files



- All the data needed for a database is stored within the data directory, commonly referred to as **PGDATA**. A common location for PGDATA is `/var/lib/pgsql/data`.

Item	Description
PG_VERSION	A file containing the major version number of PostgreSQL
base	Subdirectory containing per-database subdirectories
global	Subdirectory containing cluster-wide tables, such as <code>pg_database</code>
pg_xlog	Subdirectory containing Log files
postmaster.opts	A file recording the command-line options the server was last started with
postmaster.pid	A lock file recording the current server PID and shared memory segment ID (not present after server shutdown)

Psql command



- You can list available databases using `\l`, as follows:

```
postgres-# \l
```

- to connect/select a desired database, use `\c`

```
postgres=# \c testdb;
```

- You are now logged into PostgreSQL testdb and ready to execute your commands inside testdb.
- To exit from the database, you can use `\q`.

Create Database



Using CREATE DATABASE using SQL statement

- The basic syntax of CREATE DATABASE statement is as follows:

```
CREATE DATABASE dbname;
```

- Following is a simple example,

```
postgres=# CREATE DATABASE testdb;
```

Create Database



Using `createdb` Command

- The syntax for *createdb* is as shown below:

```
createdb [option...] [dbname [description]]
```

Parameter	Description
dbname	The name of a database to create.
description	Specifies a comment to be associated with the newly created database.
options	command-line arguments, which <code>createdb</code> accepts.

Create Database



Options

Option	Description
-e	Shows the commands being sent to the server.
-V	Print the app version and exit.
--help	Show help about dropdb command-line arguments, and exit.
-h host	Specifies the host name of the machine on which the server is running.
-p port	Specifies the TCP port on which the server is listening for connections.
-U username	User name to connect as.

Create Database



- As Example:

```
createdb -h localhost -p 5432 -U postgres testdb
```

```
password *****
```

- Above command will prompt you for password of the PostgreSQL admin user which is **postgres** by default so provide password and proceed to create your new database.

Create Database



- `CREATE DATABASE` actually works by copying an existing database. By default, it copies the standard system database named **template1**.
- If you add objects to **template1**, these objects will be copied into subsequently created user databases.

Report Time!

?

How do you create /
drop template?

Create Database



- There is a second standard system database named `template0`. This database contains the same data as the initial contents of `template1`, that is, only the standard objects predefined by your version of PostgreSQL.
- To create database using `template0`:

```
CREATE DATABASE dbname TEMPLATE template0;
```


Drop Database



Using Drop DATABASE using SQL statement

- The basic syntax of DROP DATABASE statement is as follows:

```
DROP DATABASE [ IF EXISTS ] dbname;
```

- Following is a simple example,

```
postgres=# DROP DATABASE testdb;
```

Drop Database



- We **cannot** drop a database that has any open connections, including our own connection from *psql* or *pgAdmin III*.
- We must switch to another database or *template1* if we want to delete the database we are currently connected to.

Drop Database



Using dropdb Command

The syntax for *dropdb* is as shown below:

- `dropdb [option...] dbname`

Parameter	Description
dbname	The name of a database to be deleted.
option	command-line arguments, which dropdb accepts.

Drop Database



Options

Option	Description
-e	Shows the commands being sent to the server.
-i	Issues a verification prompt before doing anything destructive.
-V	Print the dropdb version and exit.
--if-exists	Do not throw an error if the database does not exist.
--help	Show help about dropdb command-line arguments, and exit.
-h host	Specifies the host name of the machine on which the server is running.
-p port	Specifies the TCP port on which the server is listening for connections.
-U username	User name to connect as.

Drop Database



- As Example:

```
dropdb -h localhost -p 5432 -U postgres testdb
```

Password for user postgres: ****

- The above command drops database testdb.

Create Table



- Basic syntax of CREATE TABLE statement is as follows:

```
CREATE TABLE table_name(  
    column1 datatype,  
    column2 datatype,  
    column3 datatype,  
    . . . . .  
    columnN datatype,  
    PRIMARY KEY( one or more columns )  
);
```


Create Table



- Following is an example:

```
CREATE TABLE COMPANY(  
  
    ID                INT PRIMARY KEY,  
  
    NAME              TEXT,  
  
    AGE               INT,  
  
    ADDRESS            CHAR(50),  
  
    SALARY             INT  
  
);
```

Column Data Type



- PostgreSQL has a rich set of native data types available to users.
- Users can add new types to PostgreSQL using the `CREATE TYPE` command.

Numeric data type



Name	Storage Size	Description	Range
smallint	2 bytes	small-range integer	-32768 to +32767
integer	4 bytes	typical choice for integer	-2147483648 to +2147483647
bigint	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
numeric(precision, scale)	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point, number 23.5141 has a precision of 6 and a scale of 4.
smallserial	2 bytes	small autoincrementing integer	1 to 32767
serial	4 bytes	autoincrementing integer	1 to 2147483647
bigserial	8 bytes	large autoincrementing integer	1 to 9223372036854775807

Monetary data type



- The money type stores a currency amount with a fixed fractional precision.
- Input is accepted in a variety of formats, as integer and floating-point literals, as well as typical currency formatting, such as '\$1,000.00'

Name	Storage Size	Description	Range
money	8 bytes	currency amount	-92233720368547758.08 to +92233720368547758.07

Character data types



- An attempt to store a longer string will result in an error, unless the excess characters are all spaces, in this case string will be truncated to the maximum
- If the string is shorter than the declared length, values of type `character` will be space-padded; values of type `character varying` will simply store the shorter string.

Name	Description
<code>character varying(n)</code> , <code>varchar(n)</code>	variable-length with limit
<code>character(n)</code> , <code>char(n)</code>	fixed-length, blank padded
<code>text</code>	variable unlimited length

Date/Time data type



- Valid input for the time stamp types consists of the concatenation of a date and a time, followed by an optional time zone, followed by an optional AD or BC.

Name	Storage Size	Description	Examples
timestamp [without time zone]	8 bytes	both date and time (no time zone), From 4713 BC to 294276 AD	1999-01-08 04:05:06 January 8 04:05:06 99 BC
timestamp with time zone	8 bytes	both date and time, with time zone, From 4713 BC to 294276 AD	1999-01-08 04:05:06 -8:00 January 8 04:05:06 1999 PST

Date/Time data type



Name	Description
date	date (no time of day)
time [without time zone]	time of day (no date)
time with time zone	times of day only, with time zone
interval [fields]	time interval, field can be YEAR, MONTH, DAY, HOUR, MINUTE, SECOND

Boolean data type



- Valid literal values for the "true" state are:

TRUE , 't' , 'true' , 'y' , 'yes' , 'on' or '1'

- For the "false" state, the following values can be used:

FALSE, 'f' , 'false' , 'n' , 'no' , 'off' or '0'

Name	Storage Size	Description
boolean	1 byte	state of true or false

Enumerated Types



- Enumerated (enum) types are data types that comprise a static, ordered set of values.
- Enum types are created using the CREATE TYPE command, for example:

```
CREATE TYPE mood AS ENUM ('sad', 'ok', 'happy');  
SELECT * FROM person WHERE current_mood > 'ok';
```